

CASE STUDY · SUBSCRIPTION BUSINESS

Diagnosing Customer Churn

A worked walkthrough of how an analyst would approach a rising churn problem — the questions to ask, the way to structure the analysis, and the kind of findings that actually change a business decision.

The Scenario

A mid-sized subscription service (a streaming app, a SaaS tool or a meal-kit company) has seen monthly churn climb from 4% to 7% over two quarters. Leadership wants to know why and what to do about it, before the board meeting.

Step 1 — Frame the Question Before Touching Data

The instinct is to open the data and start looking. Resist it. First, write down what "why is churn rising" actually breaks down into:

- Is churn rising across **everyone** or concentrated in one segment (plan tier, signup cohort, acquisition channel)?
- Is this a **new** problem (something changed) or a **slow leak** that's always existed and just crossed a visible threshold?
- Are we losing customers at **signup**, mid-lifecycle or at **renewal** specifically?

This framing decides your whole analysis plan. Skipping it is the single biggest reason churn analyses produce vague, unactionable conclusions.

Step 2 — Segment Before You Average

A single "7% churn rate" headline number is almost always hiding the real story. The first real analytical move is a cohort breakdown: churn rate by signup month, held constant, tracked over each cohort's lifetime.

Cohort	Month 1	Month 3	Month 6	Signal
Jan signups	2%	4%	6%	Normal decay curve
Mar signups	2%	9%	14%	Something changed
May signups	2%	5%	7%	Back closer to normal

A table like this instantly does what the headline number can't: it isolates when the problem started. Here, March signups are the outlier, which turns "why is churn rising" into the much more answerable "what was different about the March cohort's experience?"

Step 3 — Find the Lever, Not Just the Symptom

With the bad cohort isolated, cross it against what you actually have control over:

- **Acquisition channel** — did March bring in more customers from a lower-intent channel (e.g. a discount-heavy ad campaign)?
- **Onboarding changes** — was there a product or pricing change that shipped around that time?
- **Support tickets** — did ticket volume or resolution time spike for that cohort specifically?

In a real investigation like this, it's common to find the answer is a combination e.g. a promotional channel brought in price-sensitive customers right as a support backlog meant their early issues went unresolved. Neither fact alone explains the churn spike; together, they do.

Step 4 — Translate Findings into a Decision

The deliverable isn't a chart, it's a recommendation leadership can act on in the board meeting. A strong closing slide does three things:

- States the **specific** finding (not "churn is up" but "churn is concentrated in the March cohort acquired via Channel X")
- Quantifies the **cost** of inaction (e.g. projected revenue impact if the pattern repeats next quarter)
- Proposes **one** clear next step (e.g. pause that channel, or fix the onboarding gap) rather than five options

What This Case Study Teaches

Churn analysis fails most often not from a lack of statistical sophistication, but from skipping the framing step and averaging away the signal. Cohort segmentation one of the simplest tools in an analyst's kit and is usually what actually finds the story.

This is a practice walkthrough for learning analytical structure — not a record of an actual client engagement. Terence Rebello · The Analyst's Toolbox